





Part II – A deep insight...

Elmar Eisemann

Computer Graphics and Visualization

Delft University of Technology

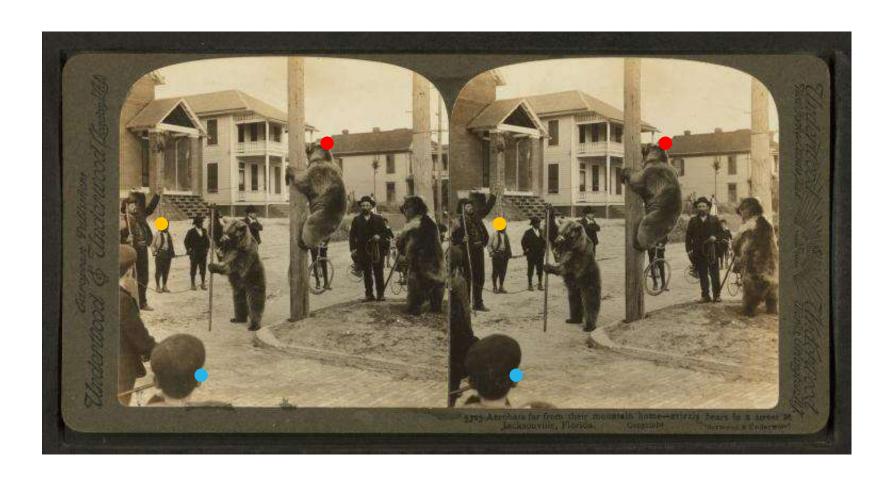






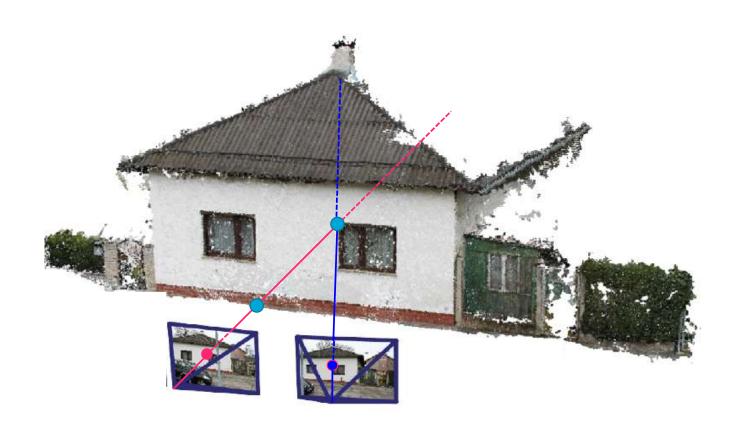
Stereograms

Stereo pairs reveal 3D information













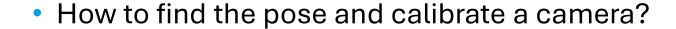
Usually not perfectly aligned

For every 3D point you need a matching pair





- How do cameras work?
 - How to model cameras mathematically?





How to calibrate from photo collections?







- How do cameras work?
 - How to model cameras mathematically?

How to find the pose and calibrate a camera?

How to reconstruct a scene?

How to calibrate from photo collections?









Producing Images in the Real World







Do you wanna build...

a camera...?

...mathematically? ©







What does this mean?

• Given 3D point M, we want a function C such that C(M) is the point's projection in the photo.







Virtual Camera Model

Projecting a scene point with the camera:

- Apply camera position (adding an offset)
- Apply rotation (matrix multiplication)
- Apply projection (non-linear scaling)

Our camera starts to become very complicated, especially if we want to calculate with it...

There has to be a better way...







Homogenous Coordinates - Definition

• N-D *projective space* Pⁿ is represented by N+1 coordinates, has no null vector, but a special equivalence relation:

Two points p, q are equal

iff (if and only if)

exists a!=0 such that p*a=q

Examples in a 2D projective space P²:

$$(2,2,2) = (3,3,3) = (4,4,4) = (\pi, \pi, \pi)$$

$$(2,2,2) \neq (3,1,3)$$

$$(0,1,0) = (0,2,0)$$

(0,0,0) not part of the space





To embed a standard vector space Rⁿ in an n-D projective space Pⁿ, we can map:

$$(x_0,x_1,...x_{n-1})$$
 in R^n to $(x_0,x_1,...x_{n-1},1)$ in P^n

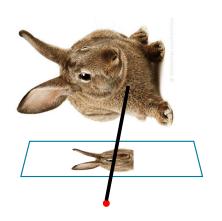
Typically, the last coordinate in a projective space is denoted with w.

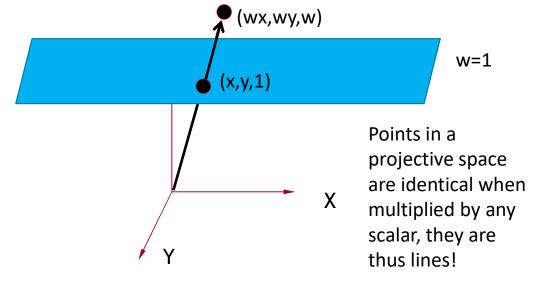




• A point (x,y) in R² embedded in a projective space corresponds to (x,y,1). All points (x,y,1) form a plane (referred to as *affine plane*)

The points in this plane correspond to points in our standard vector space R²

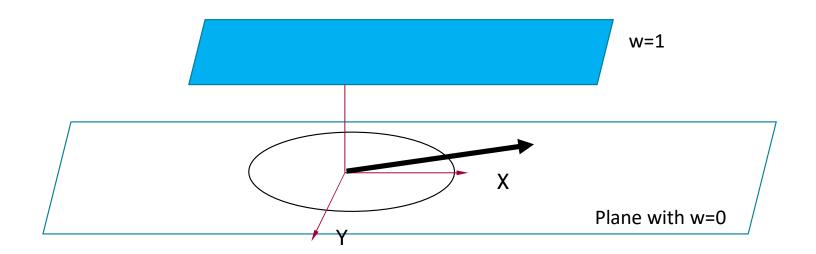








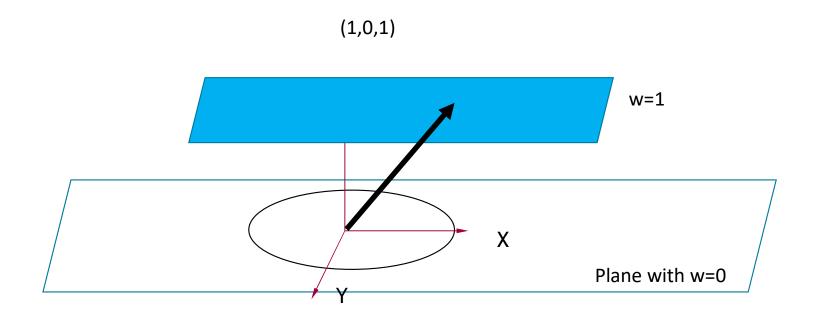
What about the points with w=0?







- What about the points with w=0?
- Let's try it out with a point (1,0,w) and we decrease w...







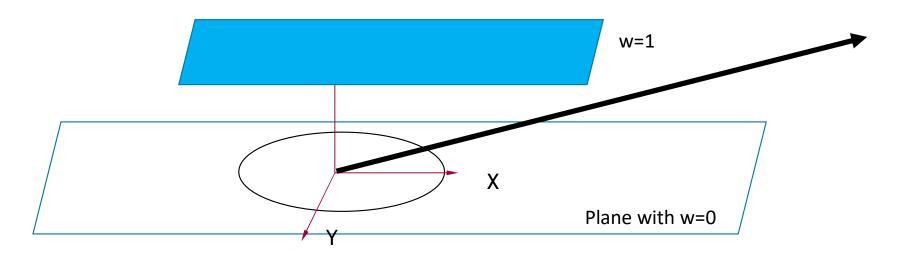
- What about the points with w=0?
- Let's try it out with a point (1,0,w) and we decrease w...





- What about the points with w=0?
- Let's try it out with a point (1,0,w) and we decrease w...

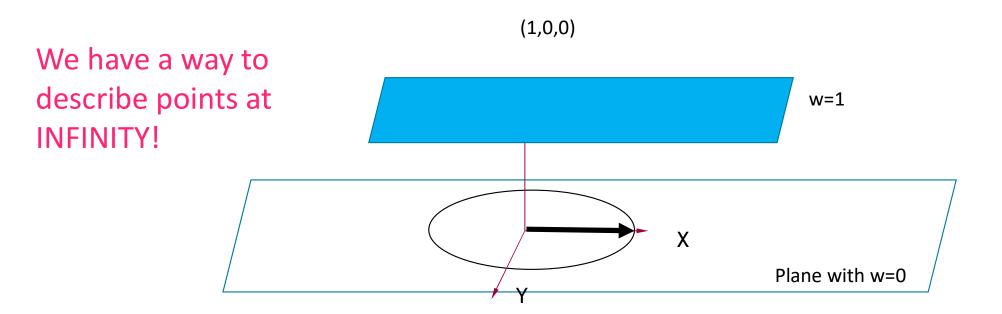
$$(1,0,0.25)=(4,0,1)$$







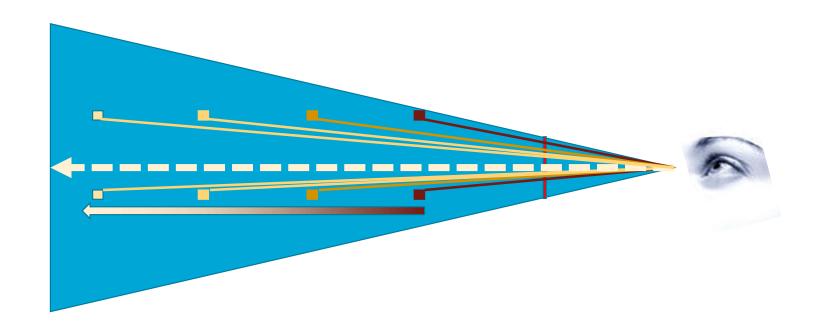
- What about the points with w=0?
- Let's try it out with a point (1,0,w) and we decrease w...







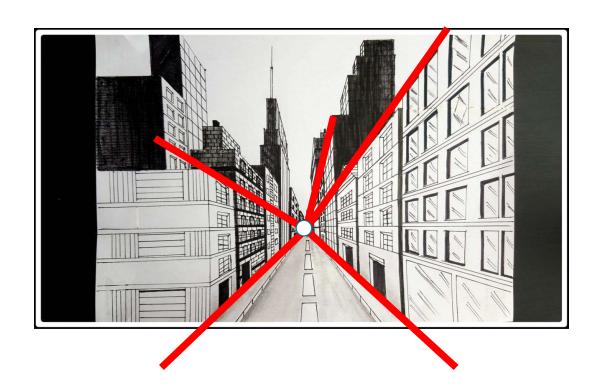
Linear Perspective







Linear Perspective







We wanted:

Easy transformations = matrices

Translations, rotations, scaling, projection should all be matrices...

Concatenating transformations means simply matrix multiplications!



6

Translations in R²

• To translate vector (x,y), we add (t_x,t_y) to obtain: $(x+t_x,y+t_y)$

In a projective space, we would like to have a matrix M, such that

$$M(x,y,1) = (x+t_x,y+t_y,1)$$





Translations in homog. coordinates

• R²

$$\begin{bmatrix} x \\ y \end{bmatrix} + \begin{bmatrix} t_x \\ t_y \end{bmatrix} = \begin{bmatrix} x + tx \\ y + ty \end{bmatrix}$$

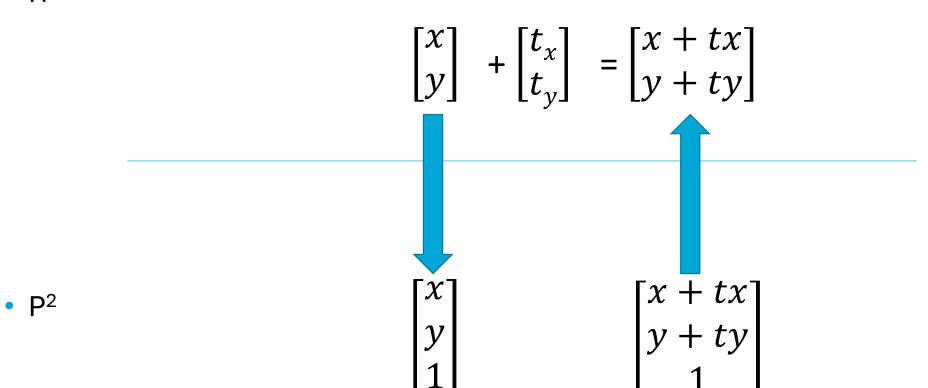
• P²





Translations in homog. coordinates

• R²





Translations in homog. coordinates

• R²

• P²

$$\begin{bmatrix} x \\ y \end{bmatrix} + \begin{bmatrix} t_x \\ t_y \end{bmatrix} = \begin{bmatrix} x + tx \\ y + ty \end{bmatrix}$$

$$\begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} x + tx \\ y + ty \\ 1 \end{bmatrix}$$



Similar stuff for rotations...





But...

• So far, we only made things more complicated!

Is there any benefit to all of this?

Yes, let's look at one example!





Rotation around point Q

- Rotation around point Q:
 - Translate Q to origin(T_O),
 - Rotate around origin (\mathbf{R}_{\odot})
 - Translate back to Q (\mathbf{T}_{-0}).

$$P'=(T_{-Q})R_{\Theta}T_{Q}P$$





Camera Model

Pinhole Perspective Projection

using matrices

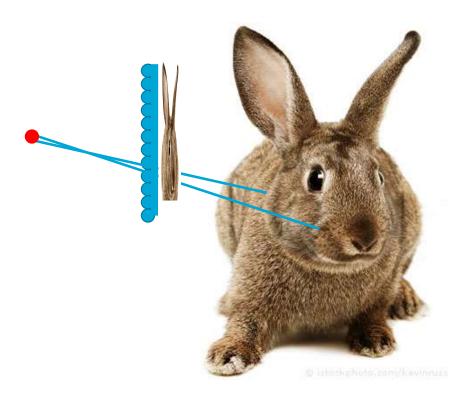






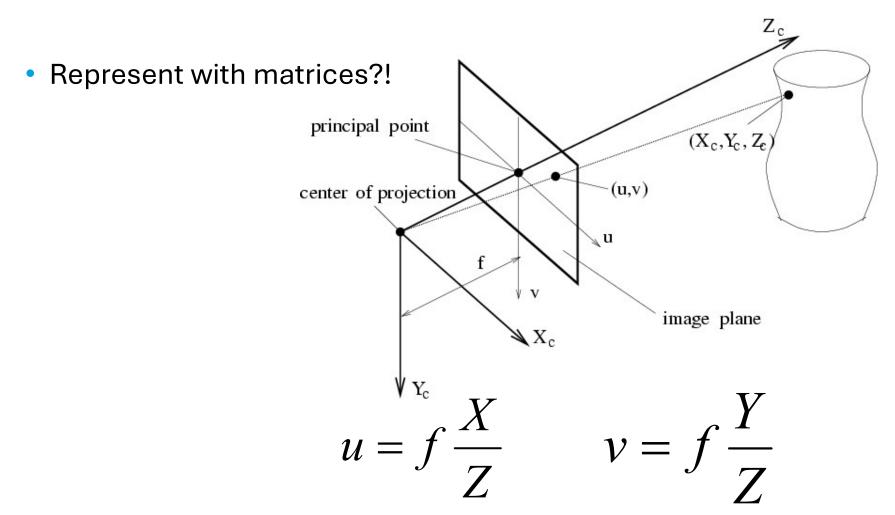
Virtual Camera

- Pinhole where real cameras have lens
- Virtual Camera Plane











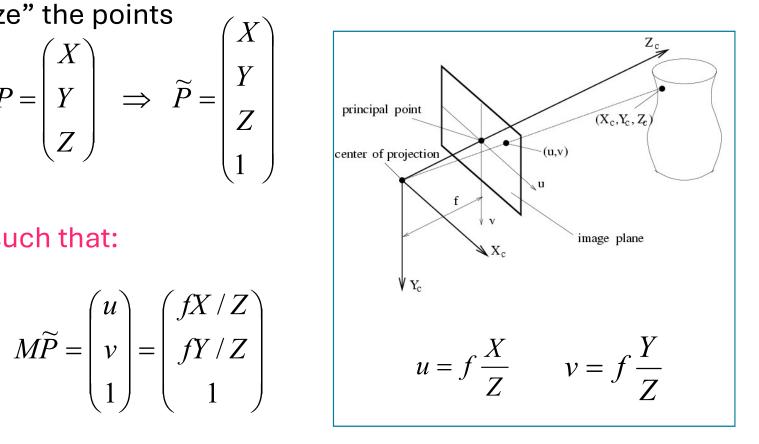


"Homogenize" the points

$$P = \begin{pmatrix} X \\ Y \\ Z \end{pmatrix} \implies \widetilde{P} = \begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix}$$
 principal point center of projection

Look for M such that:

$$M\widetilde{P} = \begin{pmatrix} u \\ v \\ 1 \end{pmatrix} = \begin{pmatrix} fX/Z \\ fY/Z \\ 1 \end{pmatrix}$$







Hint: Think projective!

$$M\widetilde{P} = \begin{pmatrix} fX/Z \\ fY/Z \\ 1 \end{pmatrix} \Rightarrow M\widetilde{P} = \begin{pmatrix} fX \\ fY \\ Z \end{pmatrix}$$





• Solution:

$$M = \begin{pmatrix} f & 0 & 0 & 0 \\ 0 & f & 0 & 0 \\ 0 & 0 & 1 & 0 \end{pmatrix}$$





Putting things together: Camera Model

Projection + Movement + Rotation

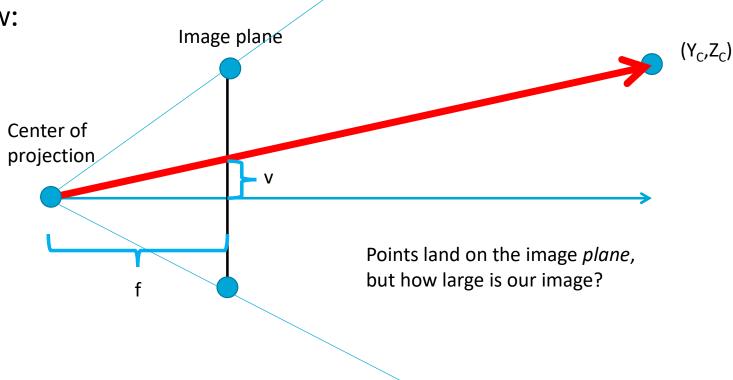
$$M = \begin{pmatrix} f & 0 & 0 & 0 \\ 0 & f & 0 & 0 \\ 0 & 0 & 1 & 0 \end{pmatrix} \begin{pmatrix} r_{00} & r_{01} & r_{02} & t_0 \\ r_{10} & r_{11} & r_{12} & t_1 \\ r_{20} & r_{21} & r_{22} & t_2 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$





Virtual Camera Model

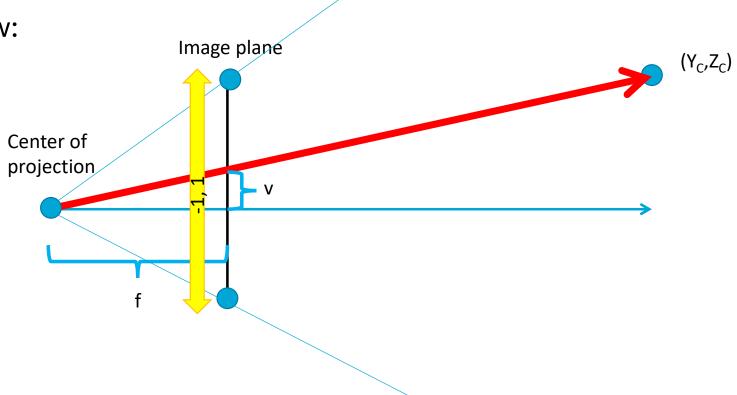
• sideview:





Virtual Camera Model

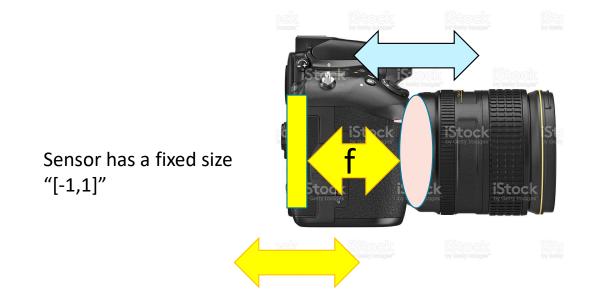
• sideview:







Real camera

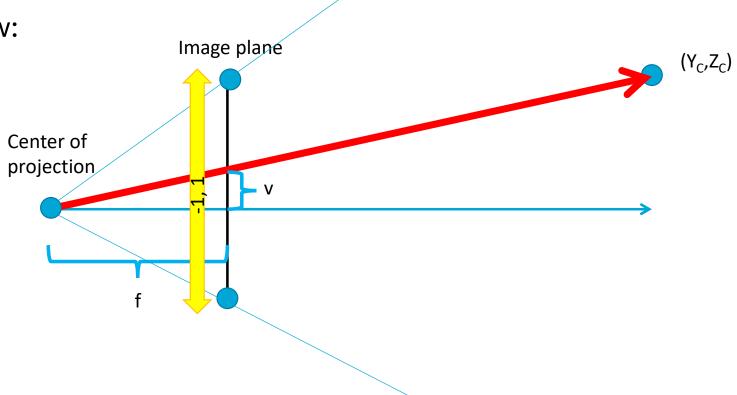






Virtual Camera Model

• sideview:

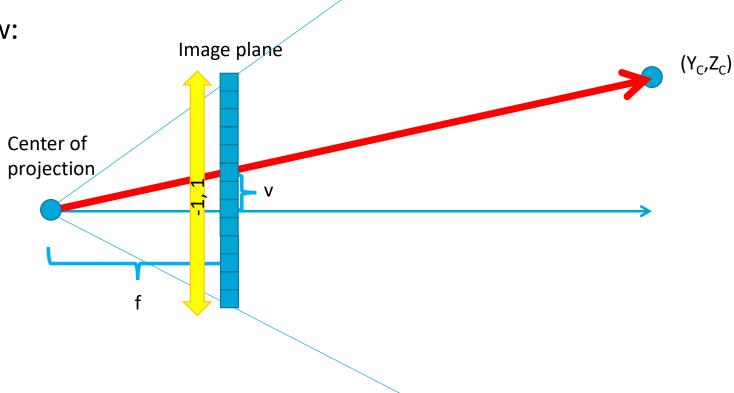






Virtual Camera Model

• sideview:

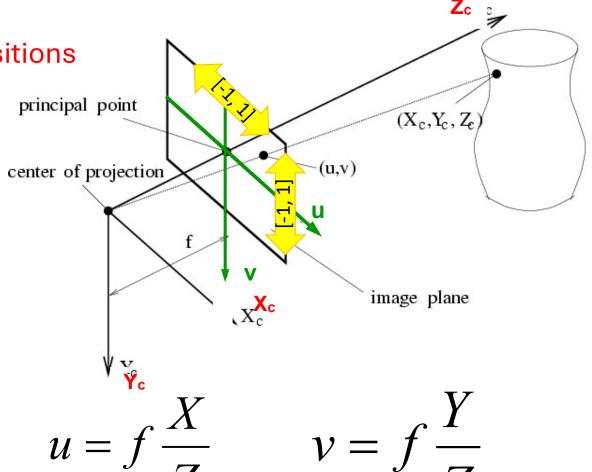




Adding Image Space to Camera Model





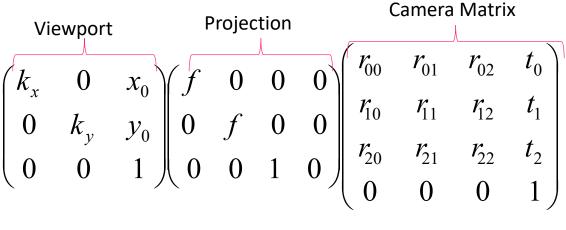


$$u = f \frac{X}{Z} \qquad v = f \frac{Y}{Z}$$





Finally:









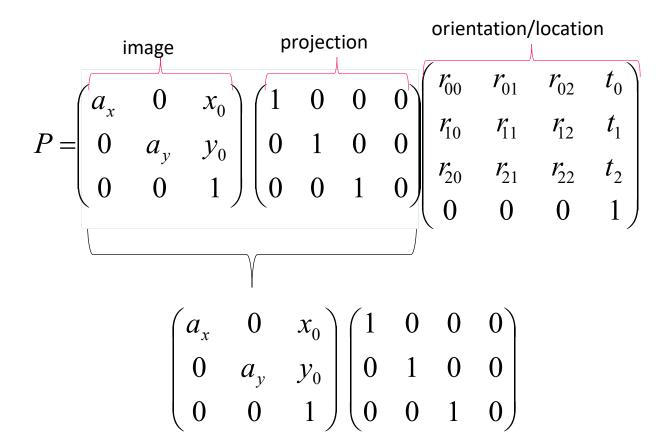
Pixel mapping

"standard camera" Deforms scene so that a "standard projection camera" can be used





Finally:







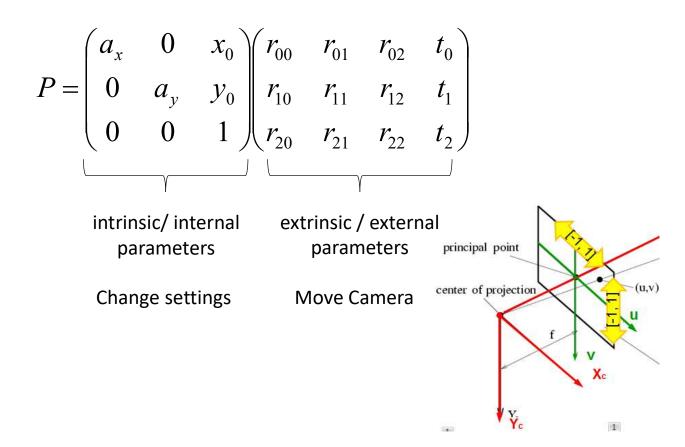
Finally:

$$P = \begin{pmatrix} a_{x} & 0 & x_{0} \\ 0 & a_{y} & y_{0} \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{pmatrix} \begin{pmatrix} r_{00} & r_{01} & r_{02} & t_{0} \\ r_{10} & r_{11} & r_{12} & t_{1} \\ r_{20} & r_{21} & r_{22} & t_{2} \\ 0 & 0 & 0 & 1 \end{pmatrix}$$
$$\begin{pmatrix} r_{00} & r_{01} & r_{02} & t_{0} \\ r_{10} & r_{11} & r_{12} & t_{1} \\ r_{20} & r_{21} & r_{22} & t_{2} \end{pmatrix}$$





Finally (notation often used in vision literature):







Finally (notation often used in vision literature):

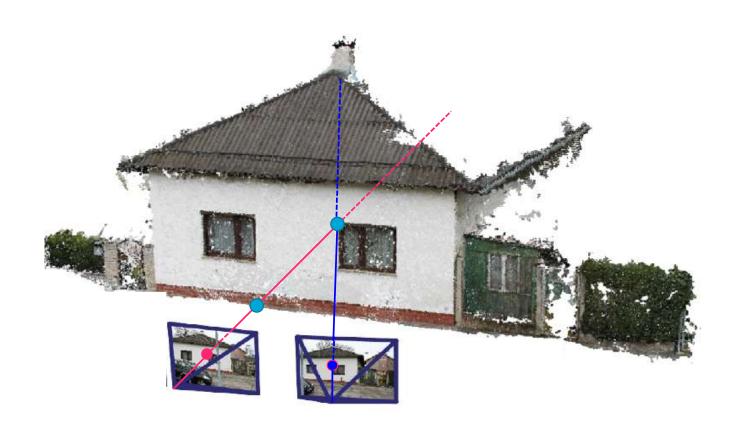
$$P = \begin{pmatrix} a_x & 0 & x_0 \\ 0 & a_y & y_0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} r_{00} & r_{01} & r_{02} & t_0 \\ r_{10} & r_{11} & r_{12} & t_1 \\ r_{20} & r_{21} & r_{22} & t_2 \end{pmatrix}$$

- Realize:
 If P is known, we can project any 3D point M to its pixel position m
- Thus, if we have a matching pixel position for cameras P1, P2, then
 we only need to solve P1(M)=m1, P2(M)=m2 for an unknown M

 Tubel



3D Reconstruction





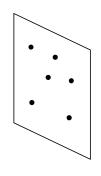


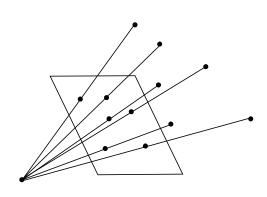
Great! But how to find the projection matrix?





Given enough known 3D points M_i and their projection m_i , we can compute camera matrix P





mi (Image Pos)

Mi (World Pos)

Hence, we determine the camera matrix P in its entirety.





- Solve for the 4x3 projection matrix P:
 - Mi (World Pos) and mi (Image Pos) are known, which means we can setup many equations of the form:

```
P Mi=mi
each such expression represents 2 equations
(2 equations= x, y position on photo)
```

12 entries in matrix P

in projective space, this means 11 unknowns; one entry can be freely chosen (e.g., P_{00} =1)

6 points in 3D with matching pixel coordinates are enough to find P





Attention:

The expression: P Mi=mi

$$\begin{pmatrix} 1 & p_{01} & p_{02} & p_{03} \\ p_{10} & p_{11} & p_{12} & p_{13} \\ p_{20} & p_{21} & p_{22} & p_{23} \end{pmatrix} \begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix} = \begin{pmatrix} kx \\ ky \\ k \end{pmatrix}$$

I see three rows?! Why 2 equations per point?

We are in a projective space!

What to do to eliminate this scalar k?





Divide first two by the last equation to get rid of the scalar k:

$$\begin{pmatrix} 1 & p_{01} & p_{02} & p_{03} \\ p_{10} & p_{11} & p_{12} & p_{13} \\ p_{20} & p_{21} & p_{22} & p_{23} \end{pmatrix} \begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix} = \begin{pmatrix} kx \\ ky \\ k \end{pmatrix}$$

$$\frac{(X + p_{01}Y + p_{02}Z + p_{03})}{(p_{10}X + p_{11}Y + p_{12}Z + p_{13})} = x$$

$$\frac{(p_{10}X + p_{11}Y + p_{12}Z + p_{13})}{(p_{20}X + p_{21}Y + p_{22}Z + p_{23})} = y$$

$$\frac{(p_{20}X + p_{21}Y + p_{22}Z + p_{23})}{(p_{20}X + p_{21}Y + p_{22}Z + p_{23})} = 1$$





$$(X + p_{01}Y + p_{02}Z + p_{03})/(p_{20}X + p_{21}Y + p_{22}Z + p_{23}) = x$$

 $(p_{10}X + p_{11}Y + p_{12}Z + p_{13})/(p_{20}X + p_{21}Y + p_{22}Z + p_{23}) = y$

Bring denominator on other side:

$$(X + p_{01}Y + p_{02}Z + p_{03}) = x(p_{20}X + p_{21}Y + p_{22}Z + p_{23}) \\ (p_{10}X + p_{11}Y + p_{12}Z + p_{13}) = y(p_{20}X + p_{21}Y + p_{22}Z + p_{23})$$

These are two linear equations with the p_{xy} being 11 unknowns to solve for.





We now solve these equations to find values for projection matrix P:

$$\begin{pmatrix} 1 & p_{01} & p_{02} & p_{03} \\ p_{10} & p_{11} & p_{12} & p_{13} \\ p_{20} & p_{21} & p_{22} & p_{23} \end{pmatrix}$$

 In consequence, P satisfies P(M_i)=m_i for all the initial points and must be the projection matrix of the camera because there is a unique solution.





Pitstop

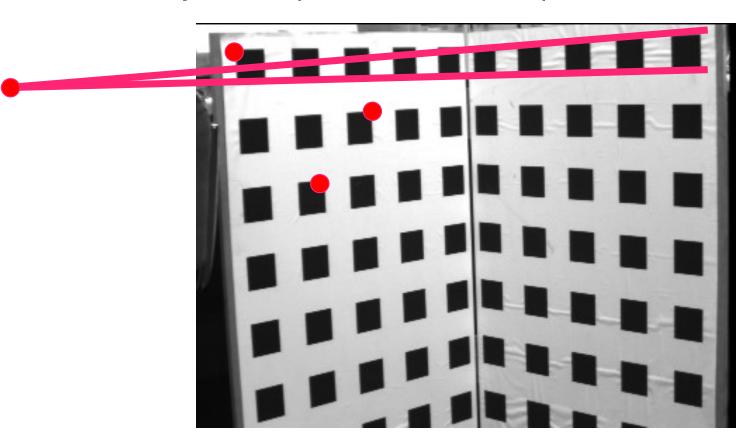
 We have seen, if we have 6 Points in 3D and their projected pixel position, we can find the camera matrix P.

How do we find 6 points and their projection?





Take a photo of something you know
 Identify known points in 3D in the photo







One way to identify points in the image:

- (i) Edge detection
- (ii) Straight line fitting to the detected edges
- (iii)Intersect lines to find images corners







- Imprecision can introduce errors:
 - world positions on the object are not perfect
 - pixel locations of projections are not perfect

- How to make things more robust?
 - Use more than 6 Points!
 - solve in a least-square sense
 (rule of thumb: 5n constraints for n unknowns)





Pitstop

- We can find the projection matrix of a camera by
 - Taking a photo of a known object
 - Choosing known points and observe their projection on the image
 - Setup a linear system with these points $P(M_i)=m_i$ and solve for the entries of P.
- What about the actual position/orientation of the camera?
 - Hard to determine from the matrix P alone...
 - We need to decompose P into the intrinsic and extrinsic camera, which makes it easier





- How do cameras work?
 - How to model cameras mathematically?





How to calibrate from photo collections?









From Camera Matrix to Intrinsic/Extrinsic Matrix

P:



We saw how to find this matrix

How can we find the actual orientation, position and intrinsic matrix?





Let's understand how orientation and camera position relate to the matrix

$$P = \begin{pmatrix} a_x & 0 & x_0 \\ 0 & a_y & y_0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} r_{00} & r_{01} & r_{02} \\ r_{10} & r_{11} & r_{12} \\ r_{20} & r_{21} & r_{22} \end{pmatrix} \begin{bmatrix} t_0 \\ t_1 \\ t_2 \end{bmatrix}$$

K (R t)

- Orientation
- Translation
- BUT: The values are not what they seem!





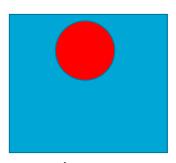
Calibration – Extracting Orientation

Why is the camera orientation not directly the rotation matrix?

$$P = \begin{pmatrix} a_x & 0 & x_0 \\ 0 & a_y & y_0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} r_{00} & r_{01} & r_{02} & t_0 \\ r_{10} & r_{11} & r_{12} & t_1 \\ r_{20} & r_{21} & r_{22} & t_2 \end{pmatrix}$$







Image





Calibration – Extracting Orientation

Why is the camera orientation not directly the rotation matrix?
 Matrix manipulates points in the world!

$$P = \begin{pmatrix} a_{x} & 0 & x_{0} \\ 0 & a_{y} & y_{0} \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} r_{00} & r_{01} & r_{02} & t_{0} \\ r_{10} & r_{11} & r_{12} & t_{1} \\ r_{20} & r_{21} & r_{22} & t_{2} \end{pmatrix}$$
Image





Calibration – Extracting Orientation

Why is the camera orientation not directly the rotation matrix?
 Matrix manipulates points in the world!

$$P = \begin{pmatrix} a_x & 0 & x_0 \\ 0 & a_y & y_0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} r_{00} & r_{01} & r_{02} & t_0 \\ r_{10} & r_{11} & r_{12} & t_1 \\ r_{20} & r_{21} & r_{22} & t_2 \end{pmatrix}$$

Camera orientation corresponds to the inverse

$$\begin{pmatrix} r_{00} & r_{01} & r_{02} \\ r_{10} & r_{11} & r_{12} \\ r_{20} & r_{21} & r_{22} \end{pmatrix}^{-1} = \begin{pmatrix} r_{00} & r_{10} & r_{20} \\ r_{01} & r_{11} & r_{21} \\ r_{02} & r_{12} & r_{22} \end{pmatrix}$$
 For rotations the inverse is the transpose





Camera Center?

Translation: Is this translation vector the position of the camera center?

$$P = \begin{pmatrix} a_x & 0 & x_0 \\ 0 & a_y & y_0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} r_{00} & r_{01} & r_{02} & t_0 \\ r_{10} & r_{11} & r_{12} & t_1 \\ r_{20} & r_{21} & r_{22} & t_2 \end{pmatrix}$$

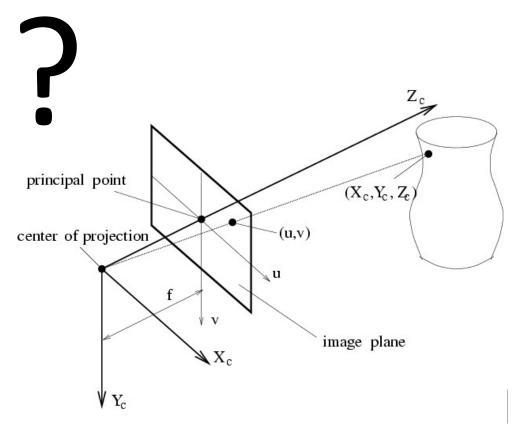
No, it is not...





If it were, then its projection should not give a finite point!

$$\begin{pmatrix} a_{x} & 0 & x_{0} \\ 0 & a_{y} & y_{0} \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} r_{00} & r_{01} & r_{02} & t_{0} \\ r_{10} & r_{11} & r_{12} & t_{1} \\ r_{20} & r_{21} & r_{22} & t_{2} \end{pmatrix} \begin{pmatrix} t_{0} \\ t_{1} \\ t_{2} \\ 1 \end{pmatrix} =$$







How can we find the camera center C from the matrix?

$$\begin{pmatrix} a_x & 0 & x_0 \\ 0 & a_y & y_0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} r_{00} & r_{01} & r_{02} & t_0 \\ r_{10} & r_{11} & r_{12} & t_1 \\ r_{20} & r_{21} & r_{22} & t_2 \end{pmatrix}$$



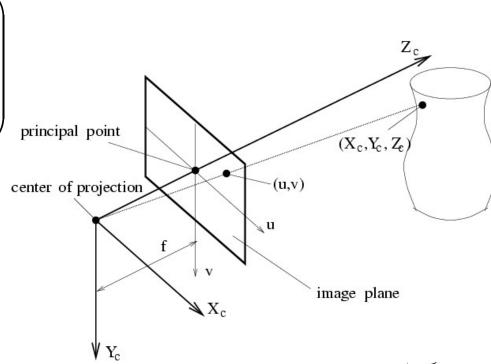


Changing C has no influence on R;
 it should be possible to derive C from t!

It can be shown that:

$$K(R \quad t) \begin{pmatrix} C_0 \\ C_1 \\ C_2 \\ 1 \end{pmatrix} = \begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix}$$

• Therefore, we have : t := -RC







Proof that t = -RC:

$$K(R - RC)\begin{pmatrix} C_0 \\ C_1 \\ C_2 \\ 1 \end{pmatrix} = K((RC) - RC) = 0$$

• Thus, C=-R-1 t, so if P=K (R t) are known, we can find the camera position





Pitstop - Calibration

- We have seen how to find the camera matrix P
- We know P = K (R^T -R^TC),
 If P is in this form, we can easily find the camera's :
- Orientation R
- Center C

• BUT: How can decompose P into matrices K and $(R^T - R^TC)$?





Camera Matrix Decomposition

Observation:

•
$$P = \begin{pmatrix} 1 & p_{01} & p_{02} & p_{03} \\ p_{10} & p_{11} & p_{12} & p_{13} \\ p_{20} & p_{21} & p_{22} & p_{23} \end{pmatrix} = K(R^T, -R^TC) = (KR^T, -KR^TC)$$





Camera Matrix Decomposition

• Observation: $P = K(R^T, -R^TC) = (KR^T, -KR^TC)$

• KR^T=
$$\begin{pmatrix} a_x & 0 & x_0 \\ 0 & a_y & y_0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} r_{00} & r_{10} & r_{20} \\ r_{01} & r_{11} & r_{21} \\ r_{02} & r_{12} & r_{22} \end{pmatrix} = \begin{pmatrix} 1 & p_{01} & p_{02} \\ p_{10} & p_{11} & p_{12} \\ p_{20} & p_{21} & p_{22} \end{pmatrix}$$

- A QR Decomposition decomposes any matrix UNIQUELY into: a rotation R and a triangular matrix Q!
- Hence, if we use the entries in P corresponding to KR^T, the decomposition has to yield K and R^T





Notation

- We understood the relationship:
- $P = K(R^{T}, -R^{T}C) = (KR^{T}, -KR^{T}C)$

- To reduce clutter, we might write:
- P = K(R', t')
- Note this is just a renaming, here: t'= -R^TC and R'= R^T



Complete Overview:

Shoot image of known object

Detect points of interest



Decompose into intrinsic/extrinsic matrix

Derive position and orientation







- How do cameras work?
 - How to model this mathematically?





How to find the pose and calibrate a camera?

How to reconstruct a scene?

How to calibrate from photo collections?

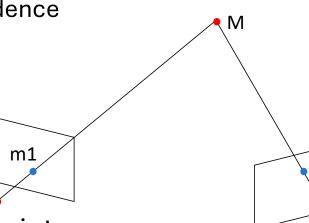




How to compute stereo

- Let's assume:
 - two calibrated cameras (known matrix)

Image point correspondence



m2

Corresponds to constraints:

$$P_1 M = m_1$$
 and $P_2 M = m_2$

$$P_2M = m_2$$

4 equations-> we can reconstruct M





How to find point correspondences?

Simple Solution:

Ask user...

Naïve solution:

Exhaustive search and compare neighborhoods around pixels

Smarter solution:

Find feature points and match those

"Smartest solution":

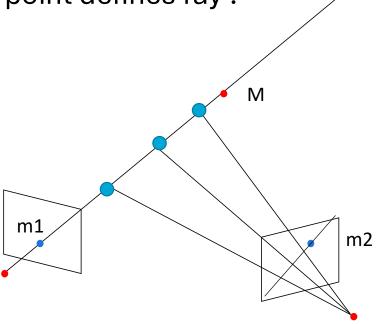
Find feature points, but only match what can geometrically work





"Smartest solution"

• Single camera: image point defines ray!

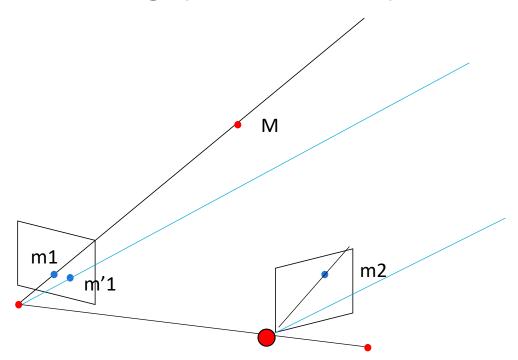






"Smartest solution"

Single camera: image point defines ray!



- Line in camera 2 is called the epipolar line
- All epipolar lines meet in an epipolar point!





"Smartest solution"

- Look along the line for matching pixel values in a neighborhood
- Image in camera 1:

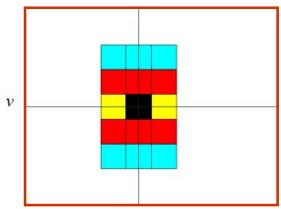
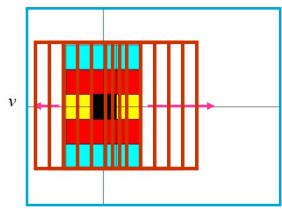


Image in camera 2:



Try all positions until the position of the most similar pixels is found.





So far so good...

- BUT:
- What happens if the cameras are **not calibrated**?





What do 2 images tell us about cameras?

- Is it possible to reconstruct the 3D configuration from 2 images?
- Probably not because we cannot obtain:
- Absolute locations
- Absolute orientation
- Scale!

"Rancor effect"



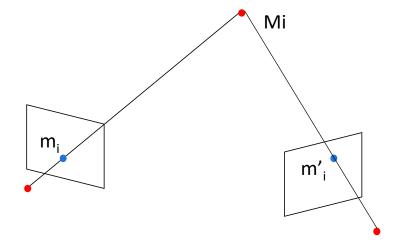




No absolute orientation or position

• Remember camera: $\begin{pmatrix} a_x & 0 & x_0 \\ 0 & a_y & y_0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} r_{00} & r_{01} & r_{02} & t_0 \\ r_{10} & r_{11} & r_{12} & t_1 \\ r_{20} & r_{21} & r_{22} & t_2 \end{pmatrix} =: K(R \quad t) =: P$

• Given $PM_i = m_i$ $P'M_i = m'_i$



, where P, P' are the projections, M_i the world and m_i, m'_i the corresponding projected points





No absolute orientation or position

• Assume we have a solution for all M_i and m_i: $P \coloneqq K(R \ t) \begin{pmatrix} X \\ Y \\ Z \\ W \end{pmatrix} = \begin{pmatrix} X \\ y \\ W \end{pmatrix}$

Affine transform leads to new solution:

S to new solution:
$$K(R \quad t) \begin{pmatrix} R' & t' \\ 0 & 1 \end{pmatrix} \begin{pmatrix} R' & t' \\ 0 & 1 \end{pmatrix}^{-1} \begin{pmatrix} X \\ Y \\ Z \\ W \end{pmatrix} = \begin{pmatrix} x \\ y \\ w \end{pmatrix}$$

$$K(R^* \quad t^*) \qquad \begin{pmatrix} X' \\ Y' \\ Z' \\ W' \end{pmatrix}$$





What if I know the intrinsic matrix?

• We have seen that we can find it by making a photo of a known object.





Finding the projection matrix

Solve for the 4x3 projection matrix P:

Mi (World Pos) and mi (Image Pos) are known, which means we can setup

many equations of the form: P Mi=mi

What if we assume that the camera defines the center of the world?

$$K(R^T - RTC) = K(Id 0) = (K 0)$$

Does knowing K help with two photos?





What if I know the intrinsic matrix?

If K is known:

$$K\begin{pmatrix} R & t \end{pmatrix} \begin{pmatrix} X \\ Y \\ Z \\ W \end{pmatrix} = \begin{pmatrix} x \\ y \\ 1 \end{pmatrix}$$

• Idea:

Multiply all pixel correspondences with K⁻¹

$$(R \quad t) \begin{pmatrix} X \\ Y \\ Z \\ W \end{pmatrix} = K^{-1} \begin{pmatrix} x \\ y \\ w \end{pmatrix} = \begin{pmatrix} x' \\ y' \\ w' \end{pmatrix}$$



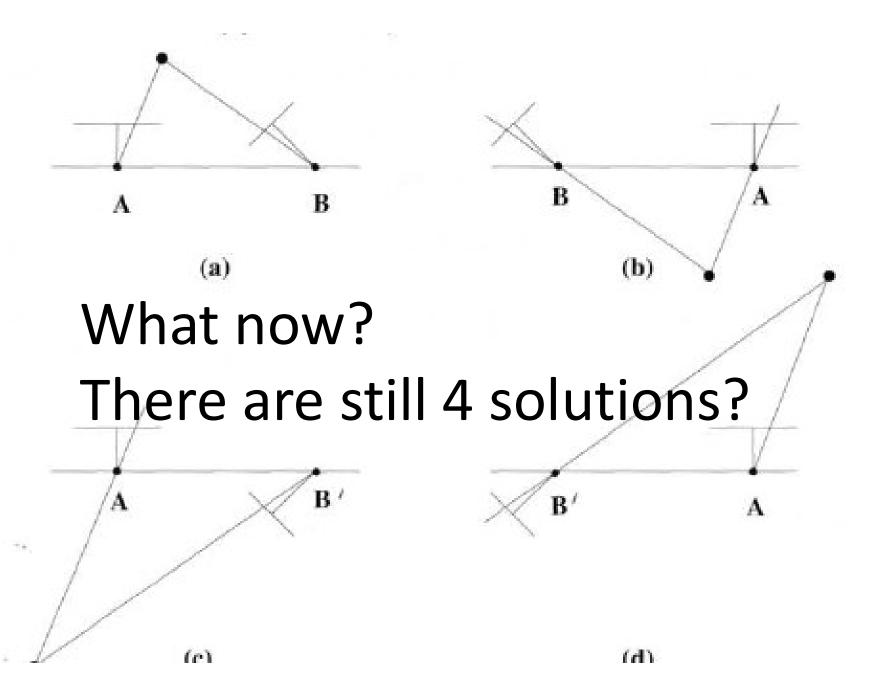


What if I know the intrinsic matrix?

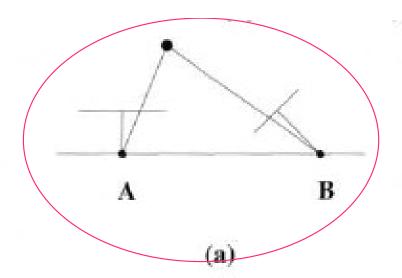
$$\begin{pmatrix} R & t \end{pmatrix} \begin{pmatrix} X \\ Y \\ Z \\ W \end{pmatrix} = \begin{pmatrix} \widetilde{x} \\ \widetilde{y} \\ \widetilde{w} \end{pmatrix}$$

- Problem becomes simpler, camera matrix is only a rotation and a translation.
- NOTE: One camera can be a reference (defining the origin and axes of the world)
 - Its rotation is identity, its origin is (0,0,0,1)^T
- In this setting, only 4 solutions exist!









What now?

There are still 4 solutions?

Only one solution has the points in front of both cameras!
All others cannot actually arise in the real world!

(d)







Pitstop

If you have found the intrinsic matrix of a camera:

- Take camera to a new place
- take two different photos
- you can reconstruct the 3D scene
 [except for scale, absolute orientation and absolute location (Rancor Effect)]
 - One of the two camera views defines your origin and world orientation





What do 2 images tell us about cameras?

For uncalibrated cameras (intrinsic matrix not known!),

Scene/Cameras are only defined up to a projective transformation.



Still better than nothing!

In fact, until [Faugeras 1993] most people even doubted that anything useful can be extracted...





No absolute orientation or position

Assume we have a solution

$$PM_i = m_i$$

 $P'M_i = m'_i$

Then PH, P'H, H-1 M_i is also a solution, for any Projective Transform H

m_i

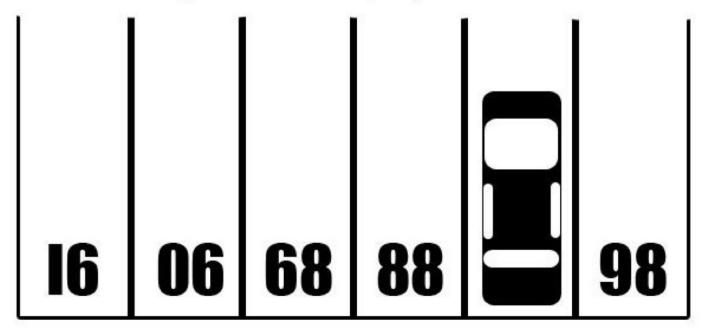
m'_i





Questions?

What is the # of the parking spot covered up by the car?

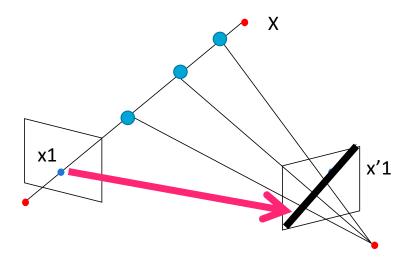






How to calibrate cameras with 2 images

Projection matrices do not well describe uncalibrated camera relations.







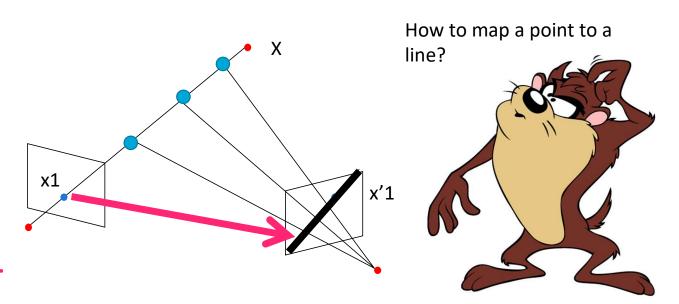
How to calibrate cameras with 2 images

Projection matrices do not well describe uncalibrated camera relations.

Uncalibrated camera relations are usually defined by the:

Fundamental Matrix F

F maps a point in one image to the epipolar line of the other.

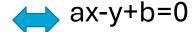


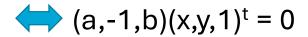


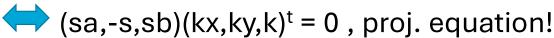
Homogeneous Coordinates

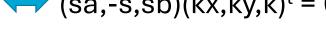
Line Point Duality in 2D

Line equation in 2D:









In other words:

Hom. points in 2D are lines!

Dot products in 2D are tests if a point is on a line!





What lines are represented by a point?

Line I through 2 points X,X'

• If X, X' are points on the line:

$$1^t X = 0$$

$$1^{t}X' = 0$$

• How can we define $!? 1 := X \wedge X'$

Point X on 2 lines l,l'

$$1'^{t} X = 0$$

$$1^t X = 0$$

$$X := 1 \wedge 1'$$





Homogenous Coordinates simplify

Because we have points at infinity,
 there are no case distinctions for parallel lines!

e.g.,
$$(a,-1,b)$$
 $(a,-1,c)=(b-c, ba-ac, 0)$

w=0 intersection at infinity!

Interesting side note: Two lines ALWAYS intersect ONCE in projective spaces!

Theorem of Bézout: Two curves of degree n, m generally have n*m intersections





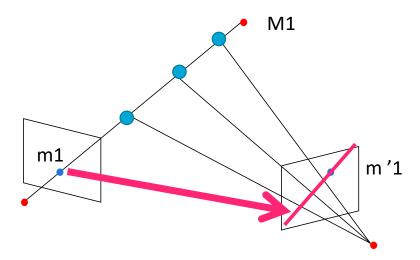
Compute F for uncalibrated cameras

Let mi, m'i be point correspondences in image plane.

Then we want F to satisfy:

$$m'_i F m_i = 0$$

(F m_i) is the epipolar line for m_i







Computing the Fundamental Matrix

Big equation system:

$m'_i F m_i = 0$



So we still solve a linear system?

8 correspondences are enough

(smart: 7, but don't worry about that) Least-square solution with more correspondences possible





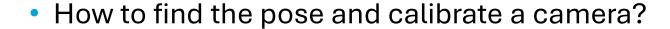
Pitstop – All the information between two images

- F encodes exactly ALL information that two images offer you!
 - It can be shown that we can build a solution with cameras directly from F
 - And a solution does imply exactly one F



Today:

- How do cameras work?
 - How to model this mathematically?



How to reconstruct a scene?

How to calibrate from photo collections?









Outlook

Outside scope of this lecture:

3 images from same camera are enough to calibrate

[Stephen J. Maybank and Olivier D. Faugeras. A theory of self-calibration of a moving camera. International Journal of Computer Vision, 8:2, 123-151, 1992]

Uses "magic" that is called trifocal tensor





Time for some music...





Summary: Guide for image collections

- Many photos come with tags...
 - If camera calibration is stored
 - -> perfect: take two cameras and reconstruct 3D
 - Camera model and settings stored
 - -> find 2 more photos with **same** camera with 3 photos you can calibrate the camera
 - Once you have known 3D points you can use those to calibrate other cameras!





Modern Solutions

- They still start with a sparse 3D point cloud reconstruction!
- Newest trend: Gaussian Splatting

- Start with 3D point reconstruction
- Replace points by Gaussians
- Optimize their position, size, and ratios to match the view from the cameras



Now we know it all...







Limitations of Stereo?









We saw:



- How to model cameras
- How to use homogenius coordinates
- How to calibrate a camera for known content
- How to calibrate with unknown content
 - How to reconstruct stereo
 - How to be more robust/faster (epipolar lines)









Thank you very much!



